

1. What is deep learning?

Deep learning is a subset of machine learning that uses neural networks with multiple layers to process and analyze complex data.

2. What are neural networks?

Neural networks are a set of algorithms designed to recognize patterns and relationships in data by mimicking the way the human brain operates.

3. What is the difference between machine learning and deep learning?

Machine learning focuses on algorithms that learn from data, while deep learning specifically uses neural networks with multiple layers for more complex tasks.

4. What is a perceptron?

A perceptron is the simplest type of neural network, consisting of a single layer that maps input features to an output using weights and an activation function.

5. What are activation functions in neural networks?

Activation functions introduce non-linearity into the neural network, allowing it to learn complex patterns. Examples include ReLU, sigmoid, and tanh.

6. What is overfitting in deep learning?

Overfitting occurs when a model learns the training data too well, including noise, and performs poorly on unseen data.

7. What is a convolutional neural network (CNN)?

A CNN is a type of deep learning model primarily used for image and video recognition tasks, employing convolutional layers to detect features like edges, textures, and shapes.

8. What is backpropagation?

Backpropagation is the algorithm used to train neural networks by calculating the gradient of the loss function and adjusting weights to minimize errors.

9. What is the role of a loss function in deep learning?

A loss function measures how well a neural network predicts the desired output. It is minimized during training to improve the model's accuracy.

10. What is the purpose of dropout in a neural network?

Dropout is a regularization technique that randomly disables neurons during training to prevent overfitting and improve generalization.

11. What is histogram equalization in image processing?

Histogram equalization is a technique used to enhance the contrast of an image by redistributing its intensity values.

12. What is thresholding in image processing?

Thresholding is a technique that converts grayscale images to binary by setting a threshold value, where pixels above the threshold become white and below become black.

13. What is edge detection in image processing?

Edge detection identifies boundaries in an image where there is a significant change in intensity, often using methods like Sobel or Canny filters.

14. What is data augmentation in the context of image processing?

Data augmentation involves artificially increasing the size of a dataset by applying random transformations (such as rotations, flips, or scaling) to existing images.

15. What are morphological operations in image processing?

Morphological operations, such as dilation and erosion, are used to process the structure of objects in an image based on their shapes.

16. What is the CIFAR-10 dataset?

CIFAR-10 is a popular dataset consisting of 60,000 32x32 color images in 10 classes, commonly used for image classification tasks.

17. How does a KNN classifier work?

KNN (K-Nearest Neighbors) classifies data by finding the majority class among the 'k' nearest data points in the feature space.

18. How does a 3-layer neural network classify data?

A 3-layer neural network consists of an input layer, one hidden layer, and an output layer, where it uses weights and activation functions to make predictions.

19. What is batch normalization?

Batch normalization normalizes the output of each layer to have zero mean and unit variance, speeding up training and improving model stability.

20. What is object detection in computer vision?

Object detection is the task of identifying and locating objects within an image using bounding boxes.

21. What is image segmentation?

Image segmentation involves partitioning an image into multiple segments or regions, usually to simplify the analysis or to identify objects.

22. What are some common tools for image labeling?

Tools like LabelImg, VGG Image Annotator (VIA), and RectLabel are used for manually labeling images for object detection and segmentation tasks.

23. What is the difference between object detection and segmentation?

Object detection identifies and locates objects, while segmentation divides the image into meaningful regions or pixels corresponding to different objects.

24. What is YOLO in object detection?

YOLO (You Only Look Once) is a real-time object detection algorithm that detects objects in a single forward pass through the network.

25. What is image captioning?

Image captioning is the process of generating descriptive text based on the content of an image using deep learning models.

26. type of neural network is typically used for generating image captions?

A combination of a convolutional neural network (CNN) for feature extraction and an RNN or LSTM for caption generation is commonly used.

27. What is a Generative Adversarial Network (GAN)?

A GAN consists of two neural networks, a generator and a discriminator, that compete with each other to generate realistic data.

28. How do GANs work?

The generator creates fake data, while the discriminator tries to distinguish between real and fake data, improving both through adversarial training.

29. What is a bi-directional LSTM?

A bi-directional LSTM processes sequences in both forward and backward directions, capturing context from both past and future inputs.

30. What is Google Colab?

Google Colab is a cloud-based platform that allows users to write and execute Python code in a Jupyter notebook environment.